

▼ Assignment 6.1

```
from keras import layers
from keras import models

model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
```

```
model.add(layers.Flatten())
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(10, activation='softmax'))
```

```
model.summary()
```

Model: "sequential"

| Layer (type) | Output Shape | Param # |
|--------------------------------|--------------------|---------|
| conv2d (Conv2D) | (None, 26, 26, 32) | 320 |
| max_pooling2d (MaxPooling2D) | (None, 13, 13, 32) | 0 |
| conv2d_1 (Conv2D) | (None, 11, 11, 64) | 18496 |
| max_pooling2d_1 (MaxPooling2D) | (None, 5, 5, 64) | 0 |
| conv2d_2 (Conv2D) | (None, 3, 3, 64) | 36928 |
| flatten (Flatten) | (None, 576) | 0 |
| dense (Dense) | (None, 64) | 36928 |
| dense_1 (Dense) | (None, 10) | 650 |
| Total params: 93,322 | | |
| Trainable params: 93,322 | | |
| Non-trainable params: 0 | | |

```
from keras.datasets import mnist
from tensorflow.keras.utils import to_categorical

(train_images, train_labels), (test_images, test_labels) = mnist.load_data()
train_images = train_images.reshape((60000, 28, 28, 1))
train_images = train_images.astype('float32') / 255
```

```
test_images = test_images.reshape((10000, 28, 28, 1))
test_images = test_images.astype('float32') / 255

train_labels = to_categorical(train_labels)
test_labels = to_categorical(test_labels)

model.compile(optimizer='rmsprop',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
model.fit(train_images, train_labels, epochs=5, batch_size=64)
```

```
Epoch 1/5
938/938 [=====] - 56s 58ms/step - loss: 0.1462 - accuracy: 0.0541
Epoch 2/5
938/938 [=====] - 55s 58ms/step - loss: 0.0440 - accuracy: 0.1543
Epoch 3/5
938/938 [=====] - 55s 59ms/step - loss: 0.0319 - accuracy: 0.1875
Epoch 4/5
938/938 [=====] - 55s 59ms/step - loss: 0.0230 - accuracy: 0.2188
Epoch 5/5
938/938 [=====] - 55s 58ms/step - loss: 0.0173 - accuracy: 0.2500
<keras.callbacks.History at 0x7ff1671e75d0>
```

```
test_loss, test_acc = model.evaluate(test_images, test_labels)
```

```
313/313 [=====] - 3s 10ms/step - loss: 0.0286 - accuracy: 0.9925
```

```
test_acc
```

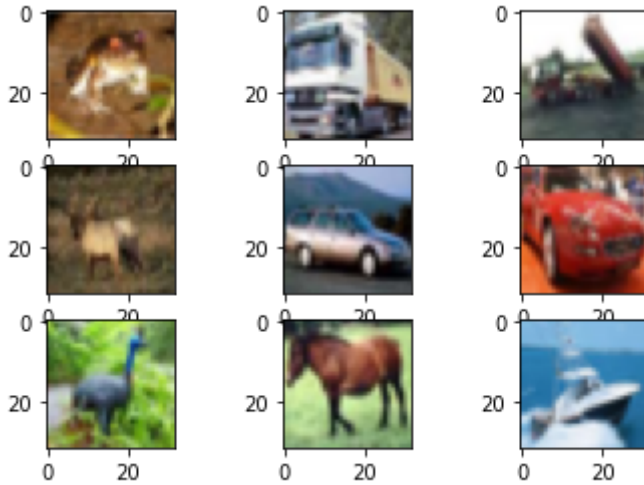
```
0.9925000071525574
```

▼ Assignment 6.2 a.

```
from matplotlib import pyplot
from keras.datasets import cifar10
# load dataset
(trainX, trainy), (testX, testy) = cifar10.load_data()
# summarize loaded dataset
print('Train: X=%s, y=%s' % (trainX.shape, trainy.shape))
print('Test: X=%s, y=%s' % (testX.shape, testy.shape))
```

```
Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
170500096/170498071 [=====] - 2s 0us/step
170508288/170498071 [=====] - 2s 0us/step
Train: X=(50000, 32, 32, 3), y=(50000, 1)
Test: X=(10000, 32, 32, 3), y=(10000, 1)
```

```
# plot first few images
for i in range(9):
    # define subplot
    pyplot.subplot(330 + 1 + i)
    # plot raw pixel data
    pyplot.imshow(trainX[i])
# show the figure
pyplot.show()
```



```
from tensorflow.keras.utils import to_categorical
# one hot encode target values
trainY = to_categorical(trainy)
testY = to_categorical(testy)
```

```
# from keras import layers
# from keras import models

# model = models.Sequential()
# model.add(layers.Conv2D(32, (3, 3), activation='relu',
#                           input_shape=(32, 32, 3)))
# model.add(layers.MaxPooling2D((2, 2)))
# model.add(layers.Conv2D(64, (3, 3), activation='relu'))
# model.add(layers.MaxPooling2D((2, 2)))
# model.add(layers.Conv2D(128, (3, 3), activation='relu'))
# model.add(layers.MaxPooling2D((2, 2)))
# model.add(layers.Conv2D(128, (3, 3), activation='relu'))
# model.add(layers.MaxPooling2D((2, 2)))
# model.add(layers.Flatten())
# model.add(layers.Dense(512, activation='relu'))
# model.add(layers.Dense(1, activation='softmax')) # sigmoid
```

```
# example of a 3-block vgg style architecture
from keras import layers
from keras import models
```

```

model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu', kernel_initializer='he_uniform'))
model.add(layers.Conv2D(32, (3, 3), activation='relu', kernel_initializer='he_uniform'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform'))
model.add(layers.Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform'))
model.add(layers.Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform'))
model.add(layers.MaxPooling2D((2, 2)))

```

```

# example output part of the model
model.add(layers.Flatten())
model.add(layers.Dense(128, activation='relu', kernel_initializer='he_uniform'))
model.add(layers.Dense(10, activation='softmax'))

```

```

# compile model
from keras import optimizers
opt = optimizers.RMSprop(lr=1e-4)
model.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy'])

```

```

/usr/local/lib/python3.7/dist-packages/tensorflow/python/keras/optimizer_v2/optimizer_v2.py:111: The `lr` argument is deprecated, use `learning_rate` instead.

```

```

# from keras import optimizers

# model.compile(loss='categorical_crossentropy', #'binary_crossentropy',
#               optimizer=optimizers.RMSprop(lr=1e-4),
#               metrics=['acc'])

```

```

# history = model.fit_generator(
#     train_generator,
#     steps_per_epoch=100,
#     epochs=30,
#     validation_data=validation_generator,
#     validation_steps=50)

```

```

history = model.fit(
    trainX, trainY,
    steps_per_epoch=100,
    epochs=20,
    validation_data=(testX, testY),
    validation_steps=50)

```

```

Epoch 1/20
100/100 [=====] - 355s 4s/step - loss: 75.4307 - accuracy: 0.0000
Epoch 2/20

```

```

100/100 [=====] - 352s 4s/step - loss: 8.6033 - accurac
Epoch 3/20
100/100 [=====] - 353s 4s/step - loss: 2.6168 - accurac
Epoch 4/20
100/100 [=====] - 351s 4s/step - loss: 1.8304 - accurac
Epoch 5/20
100/100 [=====] - 352s 4s/step - loss: 1.6397 - accurac
Epoch 6/20
100/100 [=====] - 351s 4s/step - loss: 1.5048 - accurac
Epoch 7/20
100/100 [=====] - 352s 4s/step - loss: 1.3915 - accurac
Epoch 8/20
100/100 [=====] - 351s 4s/step - loss: 1.3159 - accurac
Epoch 9/20
100/100 [=====] - 351s 4s/step - loss: 1.2402 - accurac
Epoch 10/20
100/100 [=====] - 348s 3s/step - loss: 1.1812 - accurac
Epoch 11/20
100/100 [=====] - 348s 3s/step - loss: 1.1228 - accurac
Epoch 12/20
100/100 [=====] - 348s 3s/step - loss: 1.0467 - accurac
Epoch 13/20
100/100 [=====] - 349s 3s/step - loss: 1.0122 - accurac
Epoch 14/20
100/100 [=====] - 348s 3s/step - loss: 0.9552 - accurac
Epoch 15/20
100/100 [=====] - 348s 3s/step - loss: 0.9008 - accurac
Epoch 16/20
100/100 [=====] - 349s 3s/step - loss: 0.8524 - accurac
Epoch 17/20
100/100 [=====] - 357s 4s/step - loss: 0.7881 - accurac
Epoch 18/20
100/100 [=====] - 340s 3s/step - loss: 0.7648 - accurac
Epoch 19/20
100/100 [=====] - 342s 3s/step - loss: 0.7155 - accurac
Epoch 20/20
100/100 [=====] - 341s 3s/step - loss: 0.6649 - accurac

```

```
model.save('cifar10.h5')
```

```

import matplotlib.pyplot as plt

acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']

epochs = range(1, len(acc) + 1)


plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()

```

```
plt.figure()

plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()

plt.show()
```



▼ Assignment 6.2 b.

```
# from keras.preprocessing import image
# from tensorflow.keras.preprocessing import image
from tensorflow.keras.preprocessing.image import ImageDataGenerator
# create data generator
# datagen = ImageDataGenerator(width_shift_range=0.1, height_shift_range=0.1, horizontal
datagen = ImageDataGenerator(
    rotation_range=40,
```

```

width_shift_range=0.2,
height_shift_range=0.2,
shear_range=0.2,
zoom_range=0.2,
horizontal_flip=True,
fill_mode='nearest')

```

```

# prepare iterator
it_train = datagen.flow(trainX, trainY, batch_size=64)

```

```

# example of a 3-block vgg style architecture
from keras import layers
from keras import models

model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu', kernel_initializer='he_uniform'))
model.add(layers.Conv2D(32, (3, 3), activation='relu', kernel_initializer='he_uniform'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform'))
model.add(layers.Conv2D(64, (3, 3), activation='relu', kernel_initializer='he_uniform'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform'))
model.add(layers.Conv2D(128, (3, 3), activation='relu', kernel_initializer='he_uniform'))
model.add(layers.MaxPooling2D((2, 2)))

# example output part of the model
model.add(layers.Flatten())
model.add(layers.Dropout(0.5))
model.add(layers.Dense(128, activation='relu', kernel_initializer='he_uniform'))
model.add(layers.Dense(10, activation='softmax'))

# compile model
from keras import optimizers
opt = optimizers.RMSprop(lr=1e-4)
model.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy'])

```

```

/usr/local/lib/python3.7/dist-packages/tensorflow/python/keras/optimizer_v2/optimizers.py:101: The `lr` argument is deprecated, use `learning_rate` instead.

```

```

# history = model.fit(trainX, trainY, epochs=40, validation_data=(testX, testY), batch_size=64)
# # fit model
# steps = int(trainX.shape[0] / 64)
# history = model.fit_generator(it_train, steps_per_epoch=steps, epochs=100, validation_data=(testX, testY))

```

```

history = model.fit_generator(
    it_train,
    steps_per_epoch=100,
    epochs=20,
    validation_data=(testX, testY))

```

```
validation_data=(testX, testY),
validation_steps=50)
```

```
/usr/local/lib/python3.7/dist-packages/keras/engine/training.py:1915: UserWarning:
  warnings.warn("`Model.fit_generator` is deprecated and "
```

```
Epoch 1/20
100/100 [=====] - 64s 623ms/step - loss: 132.1379 - accu
Epoch 2/20
100/100 [=====] - 62s 617ms/step - loss: 4.4708 - accur
Epoch 3/20
100/100 [=====] - 62s 620ms/step - loss: 2.4597 - accur
Epoch 4/20
100/100 [=====] - 61s 612ms/step - loss: 2.3638 - accur
Epoch 5/20
100/100 [=====] - 61s 611ms/step - loss: 2.3323 - accur
Epoch 6/20
100/100 [=====] - 61s 613ms/step - loss: 2.3304 - accur
Epoch 7/20
100/100 [=====] - 61s 613ms/step - loss: 2.3211 - accur
Epoch 8/20
100/100 [=====] - 61s 610ms/step - loss: 2.3194 - accur
Epoch 9/20
100/100 [=====] - 61s 616ms/step - loss: 2.3292 - accur
Epoch 10/20
100/100 [=====] - 61s 615ms/step - loss: 2.3183 - accur
Epoch 11/20
100/100 [=====] - 62s 619ms/step - loss: 2.3107 - accur
Epoch 12/20
100/100 [=====] - 61s 616ms/step - loss: 2.3139 - accur
Epoch 13/20
100/100 [=====] - 63s 627ms/step - loss: 2.3110 - accur
Epoch 14/20
100/100 [=====] - 63s 628ms/step - loss: 2.3095 - accur
Epoch 15/20
100/100 [=====] - 63s 634ms/step - loss: 2.3122 - accur
Epoch 16/20
100/100 [=====] - 62s 620ms/step - loss: 2.3084 - accur
Epoch 17/20
100/100 [=====] - 59s 592ms/step - loss: 2.3104 - accur
Epoch 18/20
100/100 [=====] - 57s 575ms/step - loss: 2.3150 - accur
Epoch 19/20
100/100 [=====] - 60s 600ms/step - loss: 2.3042 - accur
Epoch 20/20
100/100 [=====] - 64s 639ms/step - loss: 2.3151 - accur
```

```
import matplotlib.pyplot as plt

acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']

epochs = range(1, len(acc) + 1)

plt.plot(epochs, acc, 'bo', label='Training acc')
```





```
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()

plt.figure()

plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()

plt.show()
```



✓ 0s completed at 6:39 PM

